

Contents

ARCHITECTURE — speccheck 1.5.0	1
1. The one idea	1
2. Package layout and dependency direction	2
3. The executable flow (spec §3.2)	2
3.1 The pipeline, stage by stage	4
4. Data model	6
5. Extraction: the spec side (<code>extract.py</code>)	7
6. Extraction: the code side (<code>extract.scan_roots</code>)	8
7. Attribution (<code>attribute.py</code>)	9
8. Results mapping (<code>results.py</code>)	9
9. The status algorithm (<code>graph.py</code>)	9
10. The judge (<code>judge.py</code> , <code>judge_mock.py</code> , <code>judge_llm.py</code>)	11
10.1 Contract and validation	11
10.1a The progress indicator (R-30, C-11, K-13)	11
10.2 The mock provider	12
10.3 The LLM provider: Ollama, OpenRouter, or any OpenAI-compatible endpoint	12
11. Reporting and the write protocol (<code>report.py</code>)	13
12. The CLI (<code>cli.py</code>)	14
13. Determinism, checked three ways	15
14. The fixture and the self-check, as one artifact	15
15. Test architecture	16
15.1 Tools (<code>tools/</code>)	16
16. Things deliberately not built	17

ARCHITECTURE — speccheck 1.5.0


This document describes the system as built, module by module and data flow by data flow. It is a companion to `README.md` (how to use it), `SPEC.md` (what it must do), and `SPEC_BUILD_REPORT.md` (the evidence that it does). Spec IDs are cited inline so that every design element can be traced back to the clause that demanded it. It describes the build of `SPEC.md` v1.5: v1.2 plus the judge-stage progress indicator (v1.3), the review findings that tightened it (v1.4), and the interrupt rule made prompt (v1.5).

1. The one idea

`speccheck` answers a single question: *for every ID a specification declares, what evidence exists that the implementation realizes it?* Everything in the design follows from one boundary drawn in §0 of the spec:

The model may only ever make the news worse. Every status is computed deterministically from evidence the operator can `grep`; the judge may downgrade a status with cited evidence, never upgrade one.

That boundary splits the program into two halves that never share state:




The kernel produces every status, count, and metric. The judge, when enabled, sees one (`test case`, `ID`) edge at a time, and its only effect on the world is C-05 step 5: `PASSING` → `WEAKLY_PASSING`. Disabling it changes nothing else (I-004; property-tested by T-28).

2. Package layout and dependency direction

src/speccheck/		
__init__.py	__version__ = "1.5.0" (K-10: pyproject reads it back)	
__main__.py	python -m speccheck	
cli.py	\$5 surface; wiring; exit codes; logging; --progress; E-41; --self-check	571 lines
extract.py	C-01 grammar, SPEC.md declarations, tree walk, citations	349
attribute.py	C-03 test-case delimitation (ast) and attribution	140
results.py	C-04 JUnit parsing and the classname/join_name join	152
graph.py	C-05 status algorithm, edge selection, C-07 metrics	249
judge.py	C-06 types, validation, concurrency/budget runner, C-11 ProgressLine	338
judge_mock.py	R-22 deterministic provider	46
judge_llm.py	C-06/C-09 provider (any OpenAI-compatible endpoint)	196
judge_prompt.md	C-10 instruction text (package data, hashed into reports)	
report.py	C-07 JSON, C-08 Markdown, \$5.1 line, \$5.4 rule, \$3.1 writer	437
_selfcheck/	byte-identical copy of fixtures/target/ (package data)	

Imports flow strictly downward; no module imports cli, and nothing in the kernel imports a judge provider or an HTTP client:




Two consequences of that graph are load-bearing:

- **I-006 (network boundary)** is a property of the import layout, not of a runtime check: `httpx` is imported inside `judge_llm._httpx_post`, which only runs when `--judge llm` sends a request. `--judge none|mock` never loads an HTTP client; `--self-check` additionally installs a guard that makes `socket.socket.__init__` raise, and T-43 asserts the guard never fires.
- `graph → judge` is an import of *types only* (`JudgedVerdict`). The kernel knows the shape of a verdict so it can apply step 5; it has no idea how one is produced.

3. The executable flow (spec §3.2)

The spec draws the system once, in §3.2, as inputs feeding actors that hand artifacts to one another. This is that drawing, redrawn with the module that plays each actor:



Reading it left to right:

1. **Four inputs, three of them optional.** Only `SPEC.md` is required. Without `--src/--tests` every ID is `UNCITED`; without `junit.xml` no ID can be better than `UNVERIFIED`. Each input is read once, by exactly one actor, and never written (I-001).
2. **Two extractions that never meet.** The spec is parsed for *declarations* (what IDs exist, with statement text and retired flag); the code trees are scanned for *citations* (where each ID token occurs). The same regex recognizes the token in both, but the spec is never scanned for citations and the code is never parsed for declarations — a `~~R-07~~` in a comment is a plain citation of R-07 (E-20).
3. **Attribution gives test citations an owner.** A source citation is a (file, line); a test citation additionally names the enclosing test case (or the file-level case). That owner is what the JUnit join


and the judge both key on, which is why `TestCase` is a hashable value object (§4).

4. **The Grapher is where everything converges** — and the last point at which anything is *decided* without a model. It joins declarations to citations to outcomes, computes every status with C-05 steps 1–4, classifies undeclared and retired citations as dangling and stale, and computes the metrics. Everything on the DET artifact is reproducible from the report’s own evidence table plus `grep` and the results file (R-24).
5. **The judge is a side branch, not a stage.** With `--judge none` the deterministic result goes straight to the Reporter. Otherwise the Grapher hands over only the (`test case`, `ID`) edges of `PASSING` IDs whose test passed; the judge answers one question per edge, the kernel validates every answer, and step 5 may turn `PASSING` into `WEAKLY_PASSING`. Nothing flows back into the earlier artifacts: citations, outcomes, and counts are already fixed.
6. **Two reports, one truth.** The Reporter renders the JSON and derives the Markdown from it; the exit code and the summary line are computed from the JSON alone (I-009), so an operator holding only `speccheck.json` and the `--strict` flag can recompute what the process exited with.

The next diagram is the same flow seen as an ordered sequence of stages, with the exit paths the spec assigns to each.

3.1 The pipeline, stage by stage


`cli.execute()` runs the §3.1 stages in order. Each stage is a pure function of the previous stage’s output; there is no cache, no persistent state, and no partial output mode.



Stage timings and counts are logged at INFO (`stage=<name> ... ms=<n>`), never file contents, statements, or prompts (I-007). Exit codes are total: usage → 2, input contract → 3, any uncaught exception → 3 with a one-line message (traceback only at DEBUG), a `KeyboardInterrupt` → 3 with the message `interrupted` after the same cleanup (E-41), and everything else is the report's own `exit_code`.

4. Data model

The data model is the spec's §4 contracts, as frozen dataclasses. Nothing is mutable except the `IdRecord/TestEdge` pair that the status algorithm and step 5 fill in.



Three representational choices matter downstream:

- **IDs are normalized once, at the token boundary** (`extract.normalize_id`): R-7, R-07, R-007 become R-07; I-5 becomes I-005. Every later structure keys on the normalized string, so I-011 (injective within a family, never across families) holds by construction.
- **Test cases are hashable value objects.** `TestCase` is frozen, so it serves as the dictionary key that joins three independent producers — the attributor (which cases exist), the results mapper (which case a JUnit `<testcase>` landed on), and the judge runner (which edge a verdict belongs to) — without any ID-assignment step. The judge runner returns `dict[(TestCase, id) → JudgedVerdict]`, which is why output order never depends on completion order (K-06).
- **File-level cases are real cases with `name == ""`.** A citation outside any delimited span (module docstring, helper, fixture, an unparseable `.py`, any non-Python file) belongs to that case. They are excluded from the results join, so they can never carry an outcome and can only ever contribute UNVERIFIED (E-13) — the JSON writes `""`, the Markdown renders `(file)`.


5. Extraction: the spec side (`extract.py`)

`parse_spec` is a single pass over `SPEC.md` lines with a fence tracker (C-01, Q-008): a fence opens on a line whose first non-space characters are ````` or `~~~` and closes only on a line that is the *same* three-character marker followed by nothing but whitespace — so an info string never closes, tildes never close backticks, and an unclosed fence runs to end of file (E-04, T-05). Outside fences, two declaration forms are recognized:

Form	Rule	Example
table row	first non-space char <code>\ </code> ; cells split on unescaped <code>\ </code> outside backtick spans; the entire trimmed first cell is **ID** , ~~**ID**~~ , or **~~ID~~** ; statement = second cell	<code>\ **R-07** \ statement \ </code>
heading	# -run, then the ID token (optionally <code>~~</code> -wrapped) as the first token; statement = the rest	<code>### C-03 statement</code>

A bold ID anywhere else is not a declaration (E-31); separator rows never are. A second declaration of the same (`family`, `number`) raises `SpecError` — E-02 if both have the same retired flag, E-03 otherwise — naming both lines; there is no “first wins” (Q-006). Zero in-scope IDs is E-01. All three surface as exit 3 with no report written.

6. Extraction: the code side (`extract.scan_roots`)



Citation is literal and total: comments, strings, and code count alike (D-03), and fences are *not* excluded in source files. Exclusion by resolved path happens before size/binary filtering and is silent, which is what lets `--src` . with `--out` inside the tree stay byte-deterministic across runs (E-23, E-34, T-36). The walker sorts directory entries and the final file list, so citation order — and therefore every downstream list — is independent of filesystem order.

7. Attribution (`attribute.py`)

For `.py` files under a `--tests` root, `ast.parse` yields the test cases of C-03:

- every module-level `def test_* / async def test_*`;
- every `test*` method (underscore optional, F-102) that is a *direct* member of a top-level class recognized by name (`Test*`) or by base (`*TestCase` as a `Name` or `Attribute`).


A span runs from the first decorator line to `end_lineno` (B-03 in the build report). Methods of unrecognized or nested classes are not cases; their lines fall to the file-level case and the file gets one **undelimited tests in ...** Note (E-28). A `SyntaxError` (or the other exceptions `ast.parse` can raise on hostile input) makes the whole file file-level with a **parse fallback**: Note (E-12). Any other file type is one file-level case whose `classname` is derived exactly as for Python — path with `/→.` and the final extension dropped (F-108) — so that a JUnit result for a Go or TypeScript file can still join it by suffix.

The attributor never sees the spec: it does not know which IDs are declared. Dangling and stale classification happens in the graph, from the `SpecIndex`.

8. Results mapping (`results.py`)

`parse_junit` accepts a `<testsuites>` or `<testsuite>` root (namespaces tolerated), treats every `<testcase>` below it as one result, and derives the outcome from its children in the fixed priority **failure** > **error** > **skipped** > **passed**. A missing `name` is E-05.


The join (C-04, F-006, F-101, F-105) is two-keyed:



Everything from the first `[` is stripped for the join, so parametrized variants and duplicates land on one case whose single `outcome` is the worst of them (E-06, E-24). Unattributed results are reported, sorted, and change no status.

9. The status algorithm (`graph.py`)

`deterministic_status` is C-05 steps 1–4 verbatim; `apply_verdicts` is step 5 and is the only function in the program that reads a verdict.




`eligible_edges` selects exactly the edges the judge may see: **PASSING** after step 4 **and** the test outcome is **passed** (I-010; T-31 counts calls with a stub). The graph also classifies every citation whose ID is undeclared as *dangling* and every citation of a retired ID as *stale*; retired IDs keep empty `src/tests` so the same evidence is not listed twice (B-05).

Metrics (`compute_metrics`) use `decimal` throughout: each ratio is `Decimal(n) / Decimal(d)` under a 28-digit context, quantized to four places with `ROUND_HALF_EVEN`, or `None` when `d == 0` (I-008, Q-009). The seven in-scope statuses and the six families are always present as keys, zero or not (F-016).

10. The judge (judge.py, judge_mock.py, judge_llm.py)

10.1 Contract and validation

A provider implements one method, `judge(JudgeRequest) → Verdict`, and may raise. The kernel never trusts what comes back:



`available` (C-07 `judge_available`) is `true` unless every call that was actually made failed at the transport level; a malformed 200 counts as a call that succeeded (B-07); zero calls is vacuously available (E-36). The budget deadline is set by the first worker to issue a request and checked by every worker *before* it issues, so in-flight requests always complete and their verdicts count (Q-010; T-61).

10.1a The progress indicator (R-30, C-11, K-13)

With `--judge llm` the judge stage is the only slow part of a run, so `run_judge` accepts a `ProgressLine` (created in `cli.py` when `--progress` resolves to `on`: `auto` means `stderr` is a TTY and verbosity is not `DEBUG`). It is a context manager around the whole stage: the first draw (0/n, 0:00 elapsed, ?:?) happens on entry, before the first request; every worker calls `advance()` once its verdict is determined; a daemon thread coalesces those into at most one draw per 100 ms and adds a 1 s tick so `elapsed` and the ETA keep moving; `--exit__` draws the final state (unless an exception is propagating) and then erases. Each draw is


one write of `\r` + line + padding to the widest line so far — no terminal escapes, so a shrinking ETA leaves no residue on any console that honours a carriage return — and the erase is `\r` + spaces + `\r`. The line is written to the raw stderr stream, never through the logger, and `cli.py` defers the judge stage’s INFO lines (mode, URL, model, stage summary) until after `run_judge` returns, so the logger is silent while the line is displayed.

10.2 The mock provider

`MockJudge` reads the line-numbered `source` in the request, flags every line matching `^\s*assert\b`, containing `.assert`, or containing `pytest.raises`, and returns `ASSERTS` with those lines as evidence or `EXECUTES_ONLY` with none (R-22). It has no configuration and no I/O, which is what makes `--judge mock` byte-deterministic and usable in `--self-check`.

10.3 The LLM provider: Ollama, OpenRouter, or any OpenAI-compatible endpoint

`LlmJudge` implements C-06’s wire format against any OpenAI-compatible `chat/completions` endpoint. The reference deployment is a local **Ollama** server, which serves that shape at `http://localhost:11434/v1/chat/completions`:



Design points:

- **Exactly one request per edge, no retries** (K-06). Concurrency comes from the runner’s thread pool (`--judge-concurrency`), not from the provider.
- **K-05 is a single wall-clock deadline**, not `httpx`’s per-phase timeouts alone: the request runs on a daemon thread and the caller waits at most `SPECCKECK_JUDGE_TIMEOUT` seconds for the full body. The per-phase timeouts are set as well so the abandoned thread also ends.
- **The transport is injectable** (`LlmJudge(config, post=...)`), which is how T-33 records the exact bytes and counts concurrency, and how the CLI-level tests (T-59, T-61) stub failures, `UNKNOWN` shares, and slow calls by monkeypatching `judge_llm._httpx_post`.


- **Configuration is environment-only** (C-09): URL, model, key, optional timeout. A missing variable is a usage error naming the variable; the key's text is replaced by ******* in the one place a raw response is logged (DEBUG), and never appears anywhere else (R-23, I-007).
- **The instruction text is data, not code.** `judge_prompt.md` is package data, read through `importlib.resources`, sent verbatim as the system message, and its SHA-256 lands in the report as `judge_prompt_sha256` (R-26). T-54 asserts the file equals the C-10 block in `SPEC.md`.
- **The endpoint is a URL, nothing more.** The provider has no vendor branches: the same code talked to local Ollama for the recorded T-49 runs and to OpenRouter (`openai/gpt-4o-mini`, `--judge-concurrency 32`) for the v1.4 Phase B gate, and would talk to Anthropic's OpenAI-compatible endpoint the same way. Only the three `SPEC_CHECK_JUDGE_*` variables change (README, *LLM judge via OpenRouter*).
- **Thinking models** (Ollama returns their reasoning in `message.reasoning`) exhausted the v1.1 pin of `max_tokens: 400` before emitting the verdict; v1.2 pins 4000 (D-07, README *Thinking models and max_tokens*). The provider does not special-case this: an empty or truncated `content` is a malformed response, full stop.

11. Reporting and the write protocol (`report.py`)

`build_report` assembles the C-07 object in key order (Python dicts preserve insertion order; `judge_prompt_sha256` is inserted only for `llm`), computes `strict_judge_failure` (R-28, `unavailable` taking precedence over `unknown_rate`, Q-004) and then `exit_code` from the document itself. The Markdown is rendered *from the JSON object*, so the two can never disagree, and the §1 verdict line in the Markdown is the same `summary_line()` the CLI prints.

Ratios are Decimals but `json.dumps` cannot emit them with a fixed number of places, so a tiny sentinel type (`_Num`) is serialized through `default=` as `"\x00NUM:0.9000\x00"` and rewritten to the bare number after dumping. The result is a JSON text that is byte-determined without reference to binary floating point (Q-009): `0.9000`, never `0.9`.

The writer is the only code in the program that creates or deletes files (I-001):



`_replace` and `_make_nonce` are module-level so T-45 can inject a failure between the two renames and observe the nonce.

12. The CLI (`cli.py`)

`cli.py` owns everything the spec calls a *surface*:

- **Parsing** — an `argparse` parser whose `error()` raises `UsageError` instead of exiting, so every usage problem funnels through one `except` that logs at `ERROR` and returns 2. Manual validation covers what `argparse` cannot express with the spec’s exact messages: `--judge` values, `--verbose` levels, `--progress` modes, `--max-unknown` as a Decimal in `[0, 1]`, integer ranges, the `path` outside `--root:` rule (E-09) after a single `resolve()` per path.
- **Config** — a frozen record of resolved paths and validated options; `Action` says whether `argv` asked for a check or a self-check. T-43 records the `Config` the self-check builds.
- **Wiring** — `execute()` is the §3 pipeline above; providers are constructed by `_make_provider`, which is the seam T-31 uses to count judge calls. `_progress_enabled()` decides whether the judge stage gets a `ProgressLine` (R-30, E-39): LLM judge only, never at `DEBUG`, `auto` means `sys.stderr.isatty()`, always overrides only the TTY test. Because the indicator owns `stderr` while displayed, the judge stage’s own `INFO` lines — mode, URL and model, the stage summary — are logged only after `run_judge` returns (K-13, F-201).
- **Diagnostics** — one `logging.StreamHandler` on `stderr`, logger `speccheck`, format `%(levelname)s %(message)s`, level `ERROR` unless `--verbose`; handlers are reset on every `main()` call so the in-process test runner never accumulates them. Notes are logged at `INFO`; nothing is logged at `WARNING` (F-014). The progress indicator bypasses the logger entirely: it is written to the raw `sys.stderr` stream and erased before anything else is logged.
- **Interrupts** — `main()` catches `KeyboardInterrupt` explicitly (it is a `BaseException`, so the generic `except Exception` would let it escape with Python’s exit 130): it logs `ERROR interrupted`, a traceback only at `DEBUG`, and returns 3 — keeping K-01’s closed set of exit codes (E-41, D-16). Cleanup happens where the state lives: `write_reports` removes its temporaries and any already-renamed report on any

`BaseException`, and `ProgressLine.__exit__` erases the line, so an interrupt leaves neither a half-written report nor a half-drawn bar. Two details make the interrupt *prompt* (v1.5, E-41): `run_judge` never blocks on an untimed future wait — it polls `concurrent.futures.wait(..., timeout=0.25)`, because an untimed lock wait is not SIGINT-interruptible on macOS CPython, so a bare `pool.map()` only noticed Ctrl-C once a request happened to finish — and `LlmJudge` runs the transport in a daemon thread and waits on it in 0.25 s polls that also watch an `abort` event. On interrupt `run_judge` sets the event, cancels queued futures, and joins the pool, which returns within one poll; the abandoned HTTP requests die with the daemon threads. Measured: one SIGINT against a stub server holding requests for 15 s ends the run in 0.2 s (it was 15 s).

- **The summary line** — written as UTF-8 bytes to `stdout.buffer` when one exists, so the ASCII line is identical under a C locale or a `cp1252` stdout (R-29; T-44 runs both).
- **--self-check** — copies `_selfcheck/` to a fresh temp dir, installs the socket guard, `chdirs` there, runs `run_check(...pinned argv...)` in the same process with stdout captured, compares the two outputs with the packaged goldens byte-for-byte, removes the directory, prints `self-check: ok` (or the first mismatch), and exits 0/1. The inner run is expected to exit 1 — the fixture has planted defects — and that code is deliberately not the self-check’s result (F-107, Q-003).

13. Determinism, checked three ways

I-002 says identical inputs give byte-identical reports on any OS. The design achieves it by never letting anything volatile into the data:

Source of variation	Where it is removed
filesystem order	<code>_walk</code> sorts entries; <code>scan_roots</code> sorts files; every list in the report is sorted by a stated key (C-07)
absolute paths	<code>to_posix_relative</code> against <code>--root</code> , / separators (R-20)
dictionary / set iteration	citations are grouped in dicts but emitted sorted; verdicts are keyed by (<code>TestCase</code> , <code>id</code>)
floating point	<code>decimal</code> with fixed quantization; JSON numbers via the <code>_Num</code> sentinel
judge completion order	<code>pool.map</code> preserves request order; results are keyed, not appended
the temp-file nonce	never written into either report
the progress indicator	stderr only, drawn on a TTY (or <code>--progress always</code>) and erased; nothing about it reaches stdout or either report (T-62 compares against a <code>--progress never</code> run)
time, host, version	not emitted at all — only <code>schema_version</code>

It is then verified three ways: T-36 (same fixture, different absolute paths, `--src .`, planted temporaries → identical bytes), the checked-in goldens in `fixtures/target/golden/` (T-46), and `--self-check` from the installed package.

14. The fixture and the self-check, as one artifact

`fixtures/target/` is a 16-ID calculator project engineered so that every status and every defect class has exactly one live example — one UNCITED R, one UNTESTED C, one UNVERIFIED E, one FAILING T, one SKIPPED K, one assertion-free test (→ `WEAKLY_PASSING` under the mock judge), one dangling citation, one stale citation, one unattributed result, one file-level citation, plus a parametrized test so `results[]/param` are exercised. Its goldens were produced by the tool and hand-verified (T-24 recomputes the metrics; T-37 recomputes every status from the evidence table and the JUnit file).

`src/speccheck/_selfcheck/` is a byte-identical copy, kept in sync by `tools/sync_selfcheck.py` and guarded by T-60, and shipped in the wheel via hatch artifacts (so that root-level `.gitignore` patterns for `junit.xml/speccheck.json` cannot silently exclude it — a defect the installed-wheel check caught during the build). `golden/judge_labels.json` rides along as the hand labels for the T-49 evaluation.

15. Test architecture

The suite (`tests/`, 74 tests) mirrors the spec’s §9 groups one file per group. Each test function’s docstring starts with the T-`nn` it realizes and ends with the R/C/I/K/E ids it proves — that is what makes self-application (T-48) report every one of the 170 IDs as `PASSING`, and it is why the literal ignore-marker strings are confined to `tests/data/markers/` (F-109: a marker in a test module would make the checker ignore the test’s own citation, which happened once during the build and was caught by self-application).

Two helpers in `conftest.py` carry most of the weight: `run_cli(argv, cwd, env, tty=False)` runs `cli.main` in-process with `stdout/stderr` replaced by byte buffers, so exit codes, the summary line’s bytes, and `stderr` silence are all observable without a subprocess (`tty=True` swaps in a `TtyWrapper` whose `isatty()` is true, which is how T-63 exercises `--progress auto`; the `\r`-separated capture is what T-62 parses for the C-11 draws and erase); and `project({...})` materializes a throwaway project from a file mapping in an isolated directory per call. The judge tests use stub providers and a recording HTTP stub with a concurrency counter; nothing in the gating suite touches the network. The LLM judge itself did touch the suite in one way: the v1.4 Phase B gate (`gpt-4o-mini` via OpenRouter) judged five tests as proving their IDs only by implication (R-20, K-02, K-03, K-08, E-29), and each was strengthened to assert the behavior outright — the judge’s downgrade-only rule working as intended on this repository’s own tests.

15.1 Tools (`tools/`)

Three scripts sit beside the suite. None is collected by `pytest`; all three import `speccheck` directly and drive it in-process through `cli.run_check(argv, env, stdout)` — the same entry the self-check uses — so they exercise the real pipeline without a subprocess or a shell.

Script	Serves	What it does
<code>bench.py</code>	K-08 / T-51 (recorded, not CI-gated)	Copies the golden fixture to a temp dir and times <code>check --judge mock --strict</code> five times; then generates a 10,000-file tree (5,000 source modules in 50 packages, 5,000 test modules in 50 groups, 100 declared IDs across all six families, a matching <code>junit.xml</code>) and times it five times. Prints the medians, the machine description, and PASS/FAIL against the 2 s / 60 s bounds. The generator is deterministic, so the benchmark input is reproducible.

Script	Serves	What it does
<code>eval_judge.py</code>	T-49 (opt-in, needs Ollama)	Runs <code>check --judge llm</code> three independent times on a fresh copy of the fixture, then scores every judged edge against <code>fixtures/target/golden/judge_labels.json</code> : accuracy over non-UNKNOWN verdicts and <code>unknown_rate</code> over all judged edges, per run, no pooling. Prints model, date, <code>judge_prompt_sha256</code> , and PASS only if <i>every</i> run reaches ≥ 0.90 / ≤ 0.10 . Defaults target <code>http://localhost:11434/v1/chat/completions</code> with key <code>ollama</code> ; <code>--verbose</code> prints each edge's verdict, label, and rationale, which is how the D-07 failure mode was diagnosed.
<code>sync_selfcheck.py</code>	F-107 / T-60	Copies <code>fixtures/target/</code> into <code>src/speccheck/_selfcheck/</code> (ignoring caches and dot-files such as editor swap files); <code>--check</code> reports any drift and exits 1. Run it after editing the fixture or regenerating the goldens; T-60 fails the suite if the two trees differ.

A fourth, deliberately absent, tool: nothing regenerates the goldens automatically. They are produced by running the checker on the fixture and copying the two reports into `golden/` by hand, so that a change to the goldens is always a deliberate, reviewable diff.

16. Things deliberately not built

Per the spec's non-goals and O-2/O-3: no semantic analysis of code, no test execution, no spec-quality review, no remediation, no multi-spec runs, no daemon/watch/IDE/web surface, and no language adapters beyond Python (other languages get file-level attribution, which the join still supports by suffix). The decisions the spec's own §12 leaves open (D-01..D-16) are inherited unchanged; the build's findings against them are in `SPEC_BUILD_REPORT.md`. Two have data behind them now: D-07's `max_tokens` became SPEC v1.2, and D-16 (exit 3 on interrupt rather than the shell's 130) is what E-41 and this build implement, awaiting confirmation. Also deliberately absent: a `--exclude` for scan roots — self-application therefore cites the packaged fixture copy under `src/speccheck/_selfcheck/` as source evidence, which inflates R-01's evidence list but never a status.